

UNIT 8

Neural Fields



Johannes Brandstetter, Markus Holzleitner
Institute for Machine Learning

Copyright statement:

This material, no matter whether in printed or electronic form, may be used for personal and non-commercial educational use only. Any reproduction of this material, no matter whether as a whole or in parts, no matter whether in printed or in electronic form, requires explicit prior acceptance of the authors.

Field (physics)

- A field (in physics) is an assignment of a quantity (scalar/vector) to each point in an input domain, most commonly Euclidean space \mathbb{R}^n .
- Not to be confused with the mathematical field (“Körper” in German).
- **Neural field:** A field that is parameterized by a neural network.
- Neural fields assign scalar or vector valued quantities to each point of an input domain:

$$\begin{array}{ll}\Psi(\boldsymbol{x}) : \mathbb{R}^3 \rightarrow \mathbb{R}^3 & \text{vector field} \\ \Psi(\boldsymbol{x}) : \mathbb{R}^3 \rightarrow \mathbb{R} & \text{scalar field}\end{array}$$

Example of neural fields





- Images: $\Psi(\mathbf{x}) : \mathbb{R}^2 \rightarrow \mathbb{R}^3$ (pixel coordinates \Rightarrow rgb values)
- Geometries: $\Psi(\mathbf{x}) : \mathbb{R}^2 / \mathbb{R}^3 \rightarrow \mathbb{R}$ (volumetric coordinates \Rightarrow signed-distance function)
- Geometries: $\Psi(\mathbf{x}) : \mathbb{R}^2 / \mathbb{R}^3 \rightarrow \text{Probability}[0, 1]$ (volumetric coordinates \Rightarrow occupancy field)
- 3D scenes: $\Psi(\mathbf{x}) : \mathbb{R}^5 \rightarrow \mathbb{R}^4$ (volumetric coordinates + direction \Rightarrow density, rgb values)

Advantages

- Neural fields are:
 - compact representations (memory efficient)
 - differentiable
 - “easy” to optimize
 - easy to scale to higher dimensions
 - multimodal
- Neural fields allow for:
 - Real-time rendering
 - Memory-efficient storage
 - Resolution-invariant representation
 - ...

Example: 3D geometries

- Voxel representation
- Point representation
- Mesh representation
- Occupancy neural field (Mescheder et al)



Occupancy neural fields

- An occupancy field is defined by a continuous function that maps a 3D coordinate $\mathbf{x} \in \mathbb{R}^2/\mathbb{R}^3$ to a scalar value:
 - $\Psi(\mathbf{x}) : \mathbb{R}^2/\mathbb{R}^3 \rightarrow [0, 1]$
 - The output is the probability of the occupancy.
 - A value of 1 indicates the point is inside the object, while 0 indicates it is in empty space.
- The actual physical surface S of the object is implicitly defined as an **isosurface** (or level set) where the probability equals a specific threshold (usually 0.5).

Training occupancy neural fields

■ **Cross entropy loss:** To train the network, we use a set of sampled points x_i for which the ground truth occupancy $o_i \in \{0, 1\}$ is known. The training typically employs a Binary Cross-Entropy (BCE) loss.


- For each batch B of sampled points, we get:

$$\mathcal{L}_\theta = \sum_i^B o_i \log \Psi(x_i) + (1 - o_i) \log(1 - \Psi(x_i))$$

- Occupancy fields are often preferred for complex topologies because they are easier to train on raw point clouds or noisy data where exact distance values (needed for **signed distance fields**) are hard to calculate.

Extracting a surface

- The **Marching Cubes algorithm** is used to extract a 2D surface mesh from a 3D scalar field.
- It is the industry standard for turning **implicit representations** (like the occupancy field) into **explicit representations**.



Signed distance (neural) fields

- A Signed Distance Field (SDF) encodes not just the **state of a point** (inside or outside) but also its **proximity to the boundary**.
- In a neural SDF, a network learns a continuous mapping from a spatial coordinate to its signed distance to the nearest surface.
- A neural SDF is defined by a function $\Psi(x)$ that maps a coordinate x (2D/3D) to a scalar d : $\Psi(x) = d$
 - $d > 0$: The point is outside the object
 - $d < 0$: The point is inside the object
 - $d = 0$: The point is exactly on the surface

The Eikonal equation

- For a function to be a valid distance field, it must satisfy the Eikonal equation.
- This equation states that the **magnitude of the gradient of the distance field must be equal to 1** almost everywhere, ensuring that the field changes at a constant rate relative to Euclidean distance:
 - If $\Psi(\mathbf{x}) = d(\mathbf{x})$ is the distance to the nearest point on a surface S , then for any point \mathbf{x} not on the surface, the direction of the steepest increase in distance is directly away from the surface:

$$\|\nabla_{\mathbf{x}} \Psi(\mathbf{x})\| = 1$$

Physics-informed Eikonal loss


- The physics-informed Eikonal loss forces the network to obey the underlying geometry of a true distance field:

$$\mathcal{L}_{eik} = \mathbb{E}_x \left(\|\nabla_x \Psi(x)\| - 1 \right)^2$$

- The loss is often combined with a point-cloud loss (forcing $\Psi(x) = 0$ at the surface) to ensure the network learns both the shape and the correct distance scale.

Neural Radiance Fields

- Neural Radiance Fields (NeRFs) represent a breakthrough in 3D scene reconstruction by treating a physical space as a continuous volumetric function rather than a collection of discrete points or polygons.
- Instead of storing a “shape”, the model learns to map coordinates directly to their visual properties.




Scene representation

- Neural network replaces large N-d array:


$$\Psi(\mathbf{x}) : \mathbb{R}^5 \rightarrow \mathbb{R}^4$$

- Neural networks outputs rgb values and absorption coefficient.




Ray tracing

- Ray tracing is a computer graphics rendering technique that simulates the physical behavior of light to create highly realistic images.



Generating views with volumetric rendering

- The color of the rendered ray C is obtained by integrating along the ray.



Volume rendering equation

- The color of the rendered ray C is obtained by integrating along the ray.
- Numerically, this means summing colors $c_i = [r, g, b]^T$ along points of the ray:

$$C \approx \sum_i^N T_i \alpha_i c_i$$

- How much light is blocked along the ray:

$$T_i = \prod_{j=1}^{i-1} (1 - \alpha_j)$$


- How much light is contributed by ray segment i :

$$1 - e^{-\sigma_i \delta t_i}$$

NeRFs loss

- Given images with known camera positions.
- We sample along rays, and optimize the absorption and radiance to minimize photometric error:

$$\min_{\sigma, c} \sum_p \|\text{Render}(C^p) - C_{\text{ground truth}}^p\|^2$$



Relation to PINNs

- Usually, a neural network minimizes the error between its prediction and the ground truth (data loss)
- **Physics-informed neural networks (PINNs)** add a residual loss physics loss).
 - Physics losses use automatic differentiation to calculate the derivatives of the network's output with respect to its inputs
 - Derivatives need to satisfy the PDE
 - Example: The heat equation describes how temperature $u(x, t)$ changes over space x and time t : $\frac{\partial u}{\partial t} = \alpha \frac{\partial^2 u}{\partial x^2}$
 - We define a neural network $\hat{u}(x, t)$ that takes two inputs (x, t) and predicts the temperature u . Estimates of the terms $\frac{\partial \hat{u}}{\partial t}$ and $\frac{\partial^2 \hat{u}}{\partial x^2}$ are obtained by taking the derivative of the output $\hat{u}(x, t)$ w.r.t. spatial and temporal coordinates.
 - We minimize the residuals $\frac{\partial \hat{u}}{\partial t} - \alpha \frac{\partial^2 \hat{u}}{\partial x^2}$.

SDF is a physics-informed loss

- One of the core properties of an SDF is that they satisfy the Eikonal equation: $\|\nabla_{\mathbf{x}} \Psi(\mathbf{x})\| = \frac{1}{f(\mathbf{x})}$.
- In the special case when $f(\mathbf{x}) = 1$, $\Psi(\mathbf{x})$ becomes an SDF.
- The Eikonal residual loss therefore is

$$\mathcal{L}_{eik} = \mathbb{E}_{\mathbf{x}} \left(\|\nabla_{\mathbf{x}} \Psi(\mathbf{x})\| - 1 \right)^2.$$

- This loss is easy to optimize. No conflicting gradients!